

Beginner-Friendly Guide to TabPFN(Prepared by Siddharth Yadav)

1. What is TabPFN?

TabPFN stands for Tabular Prior-data Fine-tuned Network. It is a transformer-based model developed by Microsoft Research for tabular data classification. Unlike traditional machine learning models, TabPFN does not require training on your dataset. Instead, it performs zero-shot learning using knowledge gained from millions of synthetic tasks it was pretrained on.

2. How TabPFN Works

TabPFN treats classification as a Bayesian inference problem. During its pretraining phase, it was exposed to millions of synthetic tasks. When provided with a new dataset, TabPFN uses its internal 'prior' to make probabilistic predictions for unlabeled data points based on the given small labeled subset. It performs this via a forward pass without any model training.

3. When to Use TabPFN

TabPFN is especially useful when:

- You have a small tabular dataset (≤ 128 training samples)
- You need fast predictions without training
- You want to test hypotheses quickly
- You are working with classification (not regression)
- You want state-of-the-art performance on small tabular tasks

4. Installing TabPFN

You can install TabPFN using pip:

```
pip install tabpfn
```

Or, clone the official repository for more control:

```
git clone https://github.com/automl/TabPFN
```

```
cd TabPFN
```

```
pip install -e .
```

5. Example Code to Use TabPFN

```
import torch
from tabpfn import TabPFNClassifier
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
```

```

# Load dataset
X, y = load_iris(return_X_y=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=128, test_size=10, stratify=y)

# Preprocess data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Load TabPFN model
model = TabPFNClassifier(device='cuda' if torch.cuda.is_available() else 'cpu')

# Fit and predict
model.fit(X_train, y_train)
preds = model.predict(X_test)
print('Accuracy:', accuracy_score(y_test, preds))

```

6. Important Notes and Limitations

- TabPFN supports only classification problems (not regression).
- It works best for datasets with up to 128 training samples and 100 features.
- Categorical variables must be encoded before input.
- Running on CPU with more than 1000 samples is blocked by default for performance reasons.

7. Overcoming CPU Limitations

To use TabPFN with larger datasets on CPU, set the following environment variable before running your code:

```
export TABPFN_ALLOW_CPU_LARGE_DATASET=1
```

Or within Python:

```
import os
os.environ['TABPFN_ALLOW_CPU_LARGE_DATASET'] = '1'
```

8. Resources

- GitHub: <https://github.com/automl/TabPFN>
- Paper: <https://arxiv.org/abs/2207.01848>
- TabPFN Client API (for cloud inference): <https://github.com/PriorLabs/tabpfm-client>